3. Linear regression : predict house prices using the bottom Housing dataset

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

#load california house dataset
California = fetch_california_housing(as_frame=True)
df = california.frame
print(df.head())

#split data into features and target
x= df.drop('medHouseVal', axis=1)
Y= df['medHouseVal']

#split into training and test datasets
X_train, x_test, Y_train, Y_test = train_test_split(x, y, test_
size=0.2, random_state=42)

#create linear regression model
model = LinearRegression()
# train the model
```

```
model.fit(x_train, y_train)
# predict on test set
y_pred = model.predict(x_test)
# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f"mean Squared Error : {mse:.2f}")
print(f"R^2 Score : {r2:.2f}")
# plot Actual vs predicted
plt.figure(figsize=(8,6))
sns.scatterplot(x=y_test, y=y_pred)
plt.xlabel("Actual prices")
plt.ylabel("predicted price")
plt.title("Actual vs predicted House prices")
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], '--r')
plt.show()
```