

4 Logistic regression : classify emails as spam or not using the spam dataset

import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression

from sklearn.metrics import classification_report,

confusion_matrix, accuracy_score

from sklearn.preprocessing import StandardScaler

#1. Load the dataset

Example dataset : UCI ML Spam dataset (simulated as CSV or you can download it)

For now, let's assume you have 'spam.csv' with the last column named 'spam' (1 = spam, 0 = not spam)

df = pd.read_csv('spam.csv')

Preview the dataset

print("Dataset Preview:")

print(df.head())

#2. Prepare features and labels

X = df.drop('spam', axis=1) # Features (drop the target column)

Y = df['spam'] # Target

#3. Split the data

X_train, X_test, Y_train, Y_test = train_test_split(X, Y,

test_size=0.3, random_state=42)

#4 Feature Scaling (important for Logistic Regression)

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)

X_test_scaled = scaler.transform(X_test)

#5 Train Logistic Regression model

model = LogisticRegression()

model.fit(X_train_scaled, y_train)

#6 Predictions

y_pred = model.predict(X_test_scaled)

#7. Evaluation

print("\n Confusion Matrix:")

print(confusion_matrix(y_test, y_pred))

print("\n Accuracy Score:")

print(accuracy_score(y_test, y_pred))

~~2021~~

4th program

Dataset Preview:

	word_freq_make	word_freq_address	word_freq_all	word_freq_3d \
0	0.00	0.64	0.64	0.0
1	0.21	0.28	0.50	0.0
2	0.06	0.00	0.71	0.0
3	0.00	0.00	0.00	0.0
4	0.00	0.00	0.00	0.0

	word_freq_our	word_freq_over	word_freq_remove	word_freq_internet \
0	0.32	0.00	0.00	0.00
1	0.14	0.28	0.21	0.07
2	1.23	0.19	0.19	0.12
3	0.63	0.00	0.31	0.63
4	0.63	0.00	0.31	0.63

	word_freq_order	word_freq_mail	...	char_freq;	char_freq(\
0	0.00	0.00	...	0.00	0.000
1	0.00	0.94	...	0.00	0.132
2	0.64	0.25	...	0.01	0.143
3	0.31	0.63	...	0.00	0.137
4	0.31	0.63	...	0.00	0.135

	char_freq[_	char_freq_!	char_freq_\$	char_freq_# \
0	0.0	0.778	0.000	0.000
1	0.0	0.372	0.180	0.048
2	0.0	0.276	0.184	0.010
3	0.0	0.137	0.000	0.000

print("Confusion matrix (y-test - y) :")

print("Accuracy score :")

print("score (y-test - y) :")

3

4 0.0 0.135 0.000 0.000

	capital_run_length_average	capital_run_length_longest \
0	3.756	61
1	5.114	101
2	9.821	485
3	3.537	40
4	3.537	40

	capital_run_length_total	spam
0	278	1
1	1028	1
2	2259	1
3	191	1
4	191	1

[5 rows x 58 columns]

Confusion Matrix:

[[769 35]

[71 506]]

Classification Report:

	precision	recall	f1-score	support
0	0.92	0.96	0.94	804
1	0.94	0.88	0.91	577

accuracy		0.92	1381	
macro avg	0.93	0.92	0.92	1381
weighted avg	0.92	0.92	0.92	1381

Accuracy Score:

0.9232440260680667